A Standard of Care for Persuasive Machines

By: Jacob Iannotti

Executive brief

Consumer AI now writes, reasons aloud, and persuades at human speed and scale. Mission statements promise safety and uplift; **users experience silence, ambiguity, and uneven duty of care**. This piece names the danger (a governance gap), labels the adversary (incentives that prioritize speed over verifiable safeguards), and proposes a concrete **Standard of Care for Persuasive Machines**—with a working blueprint: equations,controls, and process scaffolding that any responsible builder can adopt.

Note to readers: This is analysis and proposal, not legal advice. No personal data is disclosed. No confidential third-party materials are included.

1) What no one is naming clearly

The threat is not "AI" in the abstract. The threat is **persuasive output without commensurate duty of care**. A large model can sound authoritative in **law, health, finance, crisis**, and social contexts where a single misstep carries outsized harm. When acknowledgment channels don't reliably confirm receipt, when triage paths are opaque, and when safety programs are not externally legible, we've crossed from "beta" into **public-risk territory**.

This is solvable, but only if we stop talking about "AI safety" as vibes and start talking about **mechanisms** we can audit.

2) A Standard of Care for Persuasive Machines (SoCPM)

Below is a pragmatic baseline any consumer-facing AI company could adopt and publish. It is technology-agnostic and immediately auditable.

A. Map

Public, plain-English description of high-risk contexts (e.g., legal, medical, financial, minors, crisis).

Model/task inventory mapped to those risks.

Clear in-product labels when a high-risk path is entered.

B. Measure

Pre-deployment evaluations for each high-risk context (benchmarks + scenario tests with failure modes).

Live "canary" evaluations in production (aggregate, privacy-preserving).

Quarterly public digest of incidents and mitigations.

C. Manage

Safety Switch / Rollback: A defined mechanism that can revert or pause models/features when harm metrics spike.

Guardrail Library: Tested controls for prohibited/redirected outputs (e.g., crisis routing, medical/legal disclaimers, de-escalation).

Approval Queue: Human-in-the-loop pathways for edge cases and escalations.

D. Govern

Lineage Ledger: A minimal record of model/data/guardrail versions that influenced a given output category (aggregate, privacy-safe), enabling blame/credit to land somewhere real.

SBOM Gate for AI: Like software SBOMs, but for model/guardrail stacks—declaring the versions and dependencies required to ship a release.

RACI for incidents: who Receives, who Acts, who Confirms, who Informs—**with SLAs**.

This is the skeleton. It fits alongside existing risk frameworks and gives the public something to read, verify, and critique

3.1 The Equations (abstracted)

At the heart is an **Equation Stack** that turns "safety" into decisions:
Context Score (Cx): likelihood the conversation is inside a high-risk domain (law/health/finance/crisis).
Vulnerability Score (V): signals that the user may be in distress or at elevated risk (purely on-platform signals; no external surveillance).
Authority Risk (Ar): how authoritative/confident the model sounds vs. its actual uncertainty.
Harm Potential (Hp): potential impact if the guidance is wrong (severity × reversibility).
Mitigation Confidence (Mc): confidence that guardrails and routing paths are functioning now (live checks).
A minimal safety decision can be described as:
SafetyDecision =
if **(Cx × Ar × Hp) – Mc × (1 – V)** exceeds a tuned threshold **T**, then **redirect/guard**; else **continue**.
Interpretation: In risky contexts with authoritative tone and high potential harm, the system **must** bias toward mitigation unless guardrails are provably strong *and* vulnerability is low. Threshold **T** is tuned per domain and must be justified with tests.

3.2 Guardrail library (selected behaviors)

Crisis Redirect: If Hp is high and V is non-zero, suppress speculative advice and present **region-appropriate crisis resources**.
Medical/Legal/Financial Boundaries: Replace prescriptive language with qualified, sources-linked guidance; require **"Confirm Understanding"** click-throughs before continuing.
Rollback Trigger: If Mc drops below a minimum (guardrail failure, stale model, incident spike), **auto-rollback** to the last known good model/guardrail set and display a banner notifying users of a safety state.
Human Escalation: When SafetyDecision triggers but the user insists on continuing, route to an **approval queue** with an explicit safety disclaimer and context snapshot.

3.3 Lineage Ledger (public-safe spec)

A privacy-preserving **ledger** that records, per release window:
model_version, guardrail_version, eval_suite_id, policy_profile, rollback_state.
Digest only (e.g., hashes, not raw data).
Append-only; exportable snapshots for external auditors.

3.4 SBOM Gate for AI (release checklist)

Shipping a release requires:
SBOM-AI.json declaring model + guardrail + eval artifacts by version/hash.
Passing **eval gates** for each declared high-risk domain.
Signed attestation by a safety owner that rollback paths are tested **this week**.

5) What companies can do this quarter

Publish a **SoCPM** page that states your Map/Measure/Manage/Govern program in plain English.
Add **receipt + case IDs** to every user safety report; publish SLA targets.
Ship a **Lineage Ledger** (aggregate, privacy-safe) and a quarterly safety digest.
Adopt an **SBOM-AI gate** before each release; include a one-paragraph rollback plan users can read.
For **health/legal/finance/crisis** flows, align copy and routing with public guidance and show users the boundary.
If you can scale a model to millions, you can scale these mechanics.

6) What regulators and press should demand

Proof that **acknowledgment and escalation** pathways exist and meet SLAs.

Public **eval summaries and failure modes** for high-risk contexts (no secrets, just substance).
Evidence that **rollback** is real (release notes + ledger entries when used).
A living **guardrail library** with domain-specific boundaries.
If a platform is persuasive at scale yet can't meet these minimums, that gulf—not AI itself—is the public danger.

8) Closing

Mission statements promise protection and uplift. **Protection is plumbing, not poetry.** The blueprint above is plumbing: thresholds, ledgers, gates, and queues that any serious team can ship. If we care about those most likely to be harmed, we will normalize this **Standard of Care for Persuasive Machines** and stop pretending that safety is mystical. It is measurable.